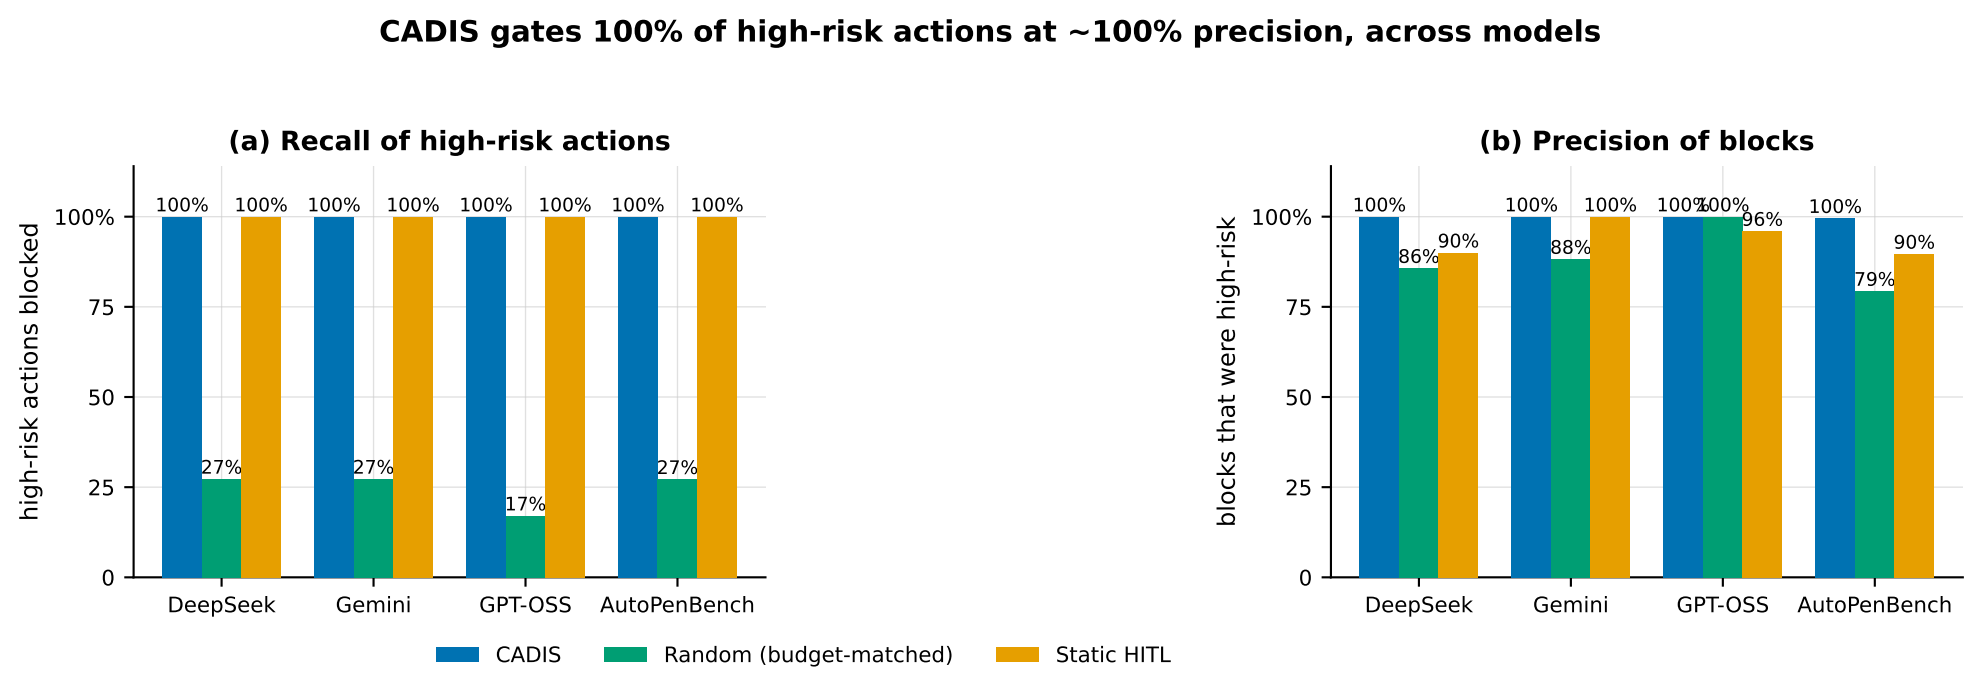
CADIS gates 100% of high-risk actions at ~100% precision, across models
(a) Recall of high-risk actions
high-risk actions blocked
100%
75
50
25
0
100%
100%
100%
100%
100%
100%
100%
100%
27%
27%
17%
27%
DeepSeek
Gemini
GPT-OSS
AutoPenBench
(b) Precision of blocks
blocks that were high-risk
100%
75
50
25
0
100%
86%
90%
100%
88%
100%
100%
100%
96%
100%
79%
90%
DeepSeek
Gemini
GPT-OSS
AutoPenBench
CADIS
Random (budget-matched)
Static HITL